

Assignment 1: Simulating a Distributed RDBMS

SimuFragDB Report

Nikunj Mahajan (IMT2023068), Gautam Kappagal (IMT2023082)
Pragya Rai (IMT2023529), Chirayu Choudhary (IMT2023004)

February 5, 2026

1 Introduction

Repository: https://github.com/Mushmat/NoSQL_Assignment1

This report documents the design, implementation, and evaluation of **SimuFragDB**, a simulated distributed database system. The objective was to implement horizontal fragmentation across multiple PostgreSQL instances and handle query routing via a Java client, gaining insights into the architectural principles of distributed systems.

2 Approach for Data Aggregation and Routing

2.1 Horizontal Fragmentation & Routing

Our implementation strictly follows the Horizontal Fragmentation strategy defined in the assignment specifications.

- **Setup:** We initialized $N = 3$ database fragments on a single physical machine.
- **Routing Logic:** We utilized the provided deterministic routing function (`Router.java`) which uses a modulo hash of the primary key:

$$\text{FragmentID} = |\text{key.hashCode}()| \pmod{\text{numFragments}}$$

- **Transactional Queries (Insert/Update/Read Profile):** For operations involving specific students (e.g., `INSERT_STUDENT`, `UPDATE_GRADE`), the request was routed explicitly to the specific fragment calculated by the router. This ensured that data for a specific entity resided consistently on one node.

2.2 Aggregation Strategy

For analytical queries such as `getAvgScoreByDept()` and `getAllStudentsWithMostCourses()`, the assignment instructions specified that “These queries must be executed on a randomly selected fragment” rather than aggregating data from all nodes.

- **Implementation:** Our client selects a random active connection from the connection pool and executes the SQL query on that single node.
- **Implication:** This design simulates a partition-local query. While efficient, it means the result set reflects only the subset of data residing on that specific shard, rather than the global dataset.

3 Challenges Faced and Resolutions

3.1 JDBC Connection Management

Challenge: Managing multiple active JDBC connections simultaneously without resource leakage.

Resolution: We implemented a `HashMap<Integer, Connection>` as a connection pool in `setupConnections()`. This allowed valid references to all fragment instances to be maintained throughout the client's lifecycle, enabling $O(1)$ retrieval of the correct connection handle for any given transaction.

3.2 Data Consistency vs. Simulation Constraints

Challenge: The primary challenge was the discrepancy between the "Expected Output" (generated from a monolithic database containing all data) and the "Actual Output" (generated from distributed fragments).

Resolution: We identified that queries like `READ_ALL` or `READ_SCORE` executed on a single random fragment would inherently return incomplete data compared to the monolithic baseline. We resolved to stick strictly to the assignment guideline of querying a "randomly selected fragment," accepting that the accuracy metric would reflect the distributed nature of the data (seeing only $1/N$ of the total records).

4 Execution Metrics

4.1 Execution Time

The total execution time for the workload as reported by the `Driver.java` was:

5529 ms

4.2 Accuracy Analysis

The evaluation compared the output of `SimuFragDB` against a baseline single-instance database.

Metric	Value
Total Query Output Lines	2462
Matching Lines	823
Accuracy Percentage	$\approx 33\%$

Table 1: Accuracy Metrics comparing Distributed vs. Monolithic execution.

Analysis of Accuracy:

The observed accuracy of approximately 33% is consistent with the system architecture. With the data horizontally fragmented across 3 nodes, any aggregation query executed on a single random node (as per instruction) has visibility into only $\approx 33\%$ of the total dataset. The remaining 67% of the data resides on the other two un-queried fragments. Thus, the low accuracy is not a defect in the code, but a verification that the data was successfully distributed and sharded across the three independent database instances.